



# Apache Spark

---

Session 12 - Machine Learning Library (MLlib) & GraphX

---

# WELCOME - KNOWBIGDATA

---

- ❑ Expert Instructors
- ❑ CloudLabs
- ❑ Lifetime access to LMS
  - ❑ Presentations
  - ❑ Class Recording
  - ❑ Assignments + Quizzes
  - ❑ Project Work
- ❑ Real Life Project
- ❑ Course Completion Certificate
- ❑ 24x7 support
- ❑ KnowBigData - Alumni
  - ❑ Jobs
  - ❑ Stay Abreast (Updated Content, Complimentary Sessions)
  - ❑ Stay Connected

# COURSE CONTENT

1	Big Data & Use Cases
2	Big Data Architectures + HDFS + YARN
3	MapReduce + NoSQL
4	HBASE + MongoDB
5	Apache Spark + Flume + Sqoop
6	Statistics & Statistical Inference
7	Understanding simple Analytics with SQL + BIRT
8	Analysing & Visualisation Data with R
9	Analysing & Visualisation Data with R (Adv)
10	Analyzing Big Data with Hive, Pig, Mahout
11	Analyzing Big Data with MLLib & SparkR
12	Advanced Visualization of Data with Tableau & D3JS

# About Instructor?

2014	<b>KnowBigData</b>	Founded
2014	<b>Amazon</b>	Built High Throughput Systems for <a href="http://Amazon.com">Amazon.com</a> site using in-house NoSql.
2012		
2012	<b>InMobi</b>	Built Recommender that churns 200 TB
2011	<b>tBits Global</b>	Founded tBits Global Built an enterprise grade Document Management System
2006	<b>D.E.Shaw</b>	Built the big data systems before the term was coined
2002	<b>IIT Roorkee</b>	Finished B.Tech.
2002		





---

# MACHINE LEARNING

---

**“Programming Computers to optimize a Performance using Example Data or Past Experience”**

- Branch of Artificial Intelligence
- Design and Development of Algorithms
- Computers Evolve Behaviour based on Empirical Data

---

# MACHINE LEARNING - TYPES

---

## **Supervised Learning**

Using Labeled training data, to create a Classifier that can predict output for unseen inputs.


## **Unsupervised Learning**

Using Unlabeled training data to create a function that can predict output.


## **Semi-Supervised Learning**

Make use of unlabeled data for training - typically a small amount of labeled data with a large amount of unlabeled data.


# MACHINE LEARNING - TYPES




Labeled Data  
(a)



Labeled and Unlabeled Data  
(b)



Supervised Learning  
(c)



Semi-Supervised Learning  
(d)



---

# MACHINE LEARNING - APPLICATIONS

---

- Recommend Friends, Dates, Products to end-user.
- Classify content into pre-defined groups.
- Find Similar content based on Object Properties.
- Identify key topics in large Collections of Text.
- Detect Anomalies within given data.
- Ranking Search Results with User Feedback Learning.
- Classifying DNA sequences.
- Sentiment Analysis/ Opinion Mining
- Computer Vision.
- Natural Language Processing,
- BioInformatics.
- Speech and HandWriting Recognition.



# MACHINE LEARNING - TOOLS

DATA SIZE	CLASSIFICATION	TOOLS
Lines Sample Data	Analysis and Visualization	Whiteboard,...
KBs - low MBs Prototype Data	Analysis and Visualization	Matlab, Octave, R, Processing,
MBs - low GBs Online Data	Analysis	NumPy, SciPy, Weka, BLAS/LAPACK
	Visualization	Flare, AmCharts, Raphael, Protovis
GBs - TBs - PBs Big Data	Analysis	Mahout, Giraph MLlib, SparkR, GraphX

---

# Machine Learning Library (MLlib)

---

Goal is to make practical machine learning scalable and easy

Consists of common learning algorithms and utilities, including:

- Classification
- Regression
- Clustering
- Collaborative filtering
- Dimensionality reduction
- Lower-level optimization primitives
- Higher-level pipeline APIs

---

# Divided into two packages

---

spark.mllib

- *contains the original API built on top of RDDs.*

spark.ml

- *provides higher-level API built on top of DataFrames for constructing ML pipelines.*



---

# spark.mllib - DataTypes

---

## Local vector

integer-typed and 0-based indices and double-typed values

```
dv2 = [1.0, 0.0, 3.0]
```

## Labeled point

a local vector, either dense or sparse, associated with a label/response

```
pos = LabeledPoint(1.0, [1.0, 0.0, 3.0])
```

## Matrices:

- Local matrix

- Distributed matrix

- RowMatrix

- IndexedRowMatrix

- CoordinateMatrix

- BlockMatrix

---

# spark.mllib - Basic Statistics

---

Summary statistics

Correlations

Stratified sampling

Hypothesis testing

Random data generation

Kernel density estimation

See <https://spark.apache.org/docs/latest/mllib-statistics.html>

---

# spark.mllib - Basic Statistics - Summary

---

```
from pyspark.mllib.stat import Statistics
sc = ... # SparkContext

mat = ... # an RDD of Vectors

# Compute column summary statistics.
summary = Statistics.colStats(mat)
print(summary.mean())
print(summary.variance())
print(summary.numNonzeros())
```



---

# MLlib - Classification and Regression

---

MLlib supports various methods:

## Binary Classification

linear SVMs, logistic regression, decision trees, random forests, gradient-boosted trees, naive Bayes

## Multiclass Classification

logistic regression, decision trees, random forests, naive Bayes

## Regression

linear least squares, Lasso, ridge regression, decision trees, random forests, gradient-boosted trees, isotonic regression

[More Details>>](#)

---

# MLlib - Collaborative Filtering

---

- Commonly used for recommender systems
- Techniques aim to fill in the missing entries of a user-item association matrix
- Supports model-based collaborative filtering,
- Users and products are described by a small set of latent factors
  - that can be used to predict missing entries.
- MLlib uses the alternating least squares (ALS) algorithm to learn these latent factors.

---

# MLlib - Collaborative Filtering - Example

---

```
from pyspark.mllib.recommendation import ALS, MatrixFactorizationModel,
Rating
# Load and parse the data
data = sc.textFile("/data/spark/mllib/als/test.data")
ratings = data.map(lambda l: l.split(',')).map(lambda l: Rating(int(l[0]), int(l[1]),
float(l[2])))

# Build the recommendation model using Alternating Least Squares
rank = 10
numIterations = 10
model = ALS.train(ratings, rank, numIterations)
```



---

# MLlib - Collaborative Filtering - Example

---


```
# Evaluate the model on training data
testdata = ratings.map(lambda p: (p[0], p[1]))
predictions = model.predictAll(testdata).map(lambda r: ((r[0], r[1]), r[2]))
ratesAndPreds = ratings.map(lambda r: ((r[0], r[1]), r[2])).join(predictions)
MSE = ratesAndPreds.map(lambda r: (r[1][0] - r[1][1])**2).mean()
print("Mean Squared Error = " + str(MSE))

# Save and load model
model.save(sc, "myModelPath")
sameModel = MatrixFactorizationModel.load(sc, "myModelPath")
```

[See More](#)

# MLlib - Clustering

- Clustering is an unsupervised learning problem
- Group subsets of entities with one another based on some notion of similarity.
- Often used for exploratory analysis



---

# MLlib supports the following models:

---

## K-means

Clusters the data points into a predefined number of clusters

## Gaussian mixture

Subgroups within overall population

## Power iteration clustering (PIC)

Clustering vertices of a graph given pairwise similarities as edge properties

## Latent Dirichlet allocation (LDA)

Infers topics from a collection of text documents

## Streaming k-means



---

# MLlib - k-means Example

---

```
from pyspark.mllib.clustering import KMeans, KMeansModel
from numpy import array
from math import sqrt

# Load and parse the data
data = sc.textFile("/data/spark/mllib/kmeans_data.txt")
parsedData = data.map(lambda line: array([float(x) for x in line.split(' ')]))

# Build the model (cluster the data)
clusters = KMeans.train(parsedData, 2, maxIterations=10, runs=10,
initializationMode="random")
```

---

# MLlib - k-means Example

---

# Evaluate clustering by computing Within Set Sum of Squared Errors

```
def error(point):
```

```
    center = clusters.centers[clusters.predict(point)]
```

```
    return sqrt(sum([x**2 for x in (point - center)]))
```

```
WSSSE = parsedData.map(lambda point: error(point)).reduce(lambda x, y: x + y)
```

```
print("Within Set Sum of Squared Error = " + str(WSSSE))
```

# Save and load model

```
clusters.save(sc, "myModelPath")
```

```
sameModel = KMeansModel.load(sc, "myModelPath")
```

---

# MLlib - Other Classes of Algorithms

---

Dimensionality reduction:

<https://spark.apache.org/docs/latest/mllib-dimensionality-reduction.html>

Feature extraction and transformation:

<https://spark.apache.org/docs/latest/mllib-feature-extraction.html>

Frequent pattern mining:

<https://spark.apache.org/docs/latest/mllib-frequent-pattern-mining.html>

Evaluation metrics:

<https://spark.apache.org/docs/latest/mllib-evaluation-metrics.html>

PMML model export:

<https://spark.apache.org/docs/latest/mllib-pmml-model-export.html>

Optimization (developer):

<https://spark.apache.org/docs/latest/mllib-optimization.html>



# GraphX

---

- For computation on graphs
- Extends the Spark RDD by introducing a new Graph abstraction
- Provides set of fundamental operators:
  - subgraph
  - joinVertices
  - aggregateMessages
- Has library of algorithms:
  - PageRank
    - If important pages link you, you are more important
  - Connected Components
    - Clusters amongst your facebook friends
  - Triangle Counting
    - Triangles passing through each vertex => measure of clustering.
- Only available in Scala yet

# GraphX - Pagerank

---

```
// Load the edges as a graph
val graph = GraphLoader.edgeListFile(sc, "graphx/data/followers.txt")
// Run PageRank
val ranks = graph.pageRank(0.0001).vertices
// Join the ranks with the usernames
val users = sc.textFile("graphx/data/users.txt").map { line =>
    val fields = line.split(",")
    (fields(0).toLong, fields(1))
}
val ranksByUsername = users.join(ranks).map {
    case (id, (username, rank)) => (username, rank)
}
// Print the result
println(ranksByUsername.collect().mkString("\n"))
```

[See more](#)



# Apache Spark

---

Thank you.

+1 419 665 3276 (US)  
+91 803 959 1464 (IN)

[reachus@knowbigdata.com](mailto:reachus@knowbigdata.com)

Subscribe to our Youtube channel for latest videos - <https://www.youtube.com/channel/UCxugRFe5wETYA7nMH6VGyEA>